

Convergence of Random Reshuffling Under the Kurdyka-Łojasiewicz Inequality

Xiao Li

with Andre Milzarek and Junwen Qiu

School of Data Science
The Chinese University of Hong Kong, Shenzhen

<https://arxiv.org/pdf/2110.04926>



Outline

Random Reshuffling (RR)

Standard Kurdyka-Łojasiewicz (KL) Analysis Framework

A New Analysis Framework and Main Results

Conclusion



Overview

Random Reshuffling (RR)

Standard Kurdyka-Łojasiewicz (KL) Analysis Framework

A New Analysis Framework and Main Results

Conclusion



Problem Statement


Optimization Problem

Consider the **finite-sum** optimization problem


$$\min_{x \in \mathbb{R}^n} f(x) := \frac{1}{N} \sum_{i=1}^N f(x, i),$$

$f(\cdot, i)$ is **smooth** and **possibly nonconvex** $\forall i \in [N] := \{1, \dots, N\}$.


Applications:



Neural Network



Supervised Learning



Matrix Optimization



Random Reshuffling (RR)

RR: Iteration t

1. Set $x_0^t = x^t$ and generate a **random permutation** σ^t of $[N]$;
2. **For** $i = 1, \dots, N$ (loop of one epoch)
 $x_i^t = x_{i-1}^t - \alpha_t \nabla f(x_{i-1}^t, \sigma_i^t);$

End

3. Set $x^{t+1} = x_N^t$.

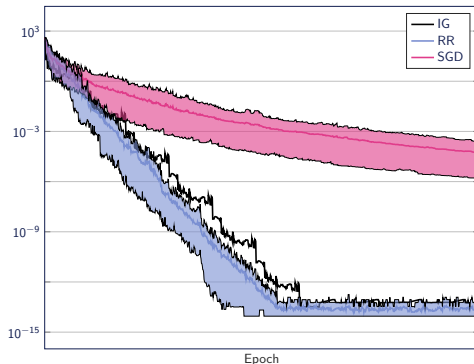
Remark:

- ▶ The step size sequence $\{\alpha_t\}$ is **diminishing**.
- ▶ **Incremental gradient (IG) method:** $\sigma^t = [N]$.
- ▶ **Stochastic gradient (SGD) method:** sampling **with-replacement**.

IG	f_1	f_2	f_3	f_4	f_5
RR	f_2	f_5	f_4	f_1	f_3
SGD	f_3	f_2	f_2	f_1	f_3



Why Random Reshuffling (RR)?



- ▶ Easy to implement due to **simple implementation**.
- ▶ Better performance than **SGD**.
- ▶ Widely used in **large-scale** learning software packages.



Different Types of Convergence:

Two main types [Absil et al. (2005)]:

- ▶ **Weak convergence:** $\lim_{t \rightarrow \infty} \|\nabla f(x^t)\| = 0$;
- ▶ **Strong (limit-point) convergence:** $\lim_{t \rightarrow \infty} x^t = x^* \in \text{crit}(f)$;

Another usually used notion of "convergence":

- ▶ **Iteration complexity:** No convergence guarantees for $\|\nabla f(x^t)\|$,

$$\min_{0 \leq t \leq T} \|\nabla f(x^t)\| \leq \frac{1}{\sqrt{T}}, \quad \text{or} \quad \frac{1}{T} \sum_{t=1}^T \|\nabla f(x^t)\| \leq \frac{1}{\sqrt{T}}.$$



Existing Theoretical Results

IG

- ▶ **Weak convergence** has been shown in [Luo and Tseng, 1994; Tseng, 1998; Solodov and Zavriev, 1998].
- ▶ **Strong convergence** at rate $\mathcal{O}(1/t)$ under **strong convexity** has been shown in [Gürbüzbalaban et al., 2019].



Existing Theoretical Results

IG

- ▶ **Weak convergence** has been shown in [Luo and Tseng, 1994; Tseng, 1998; Solodov and Zavriev, 1998].
- ▶ **Strong convergence** at rate $\mathcal{O}(1/t)$ under **strong convexity** has been shown in [Gürbüzbalaban et al., 2019].

RR

- ▶ **Strong convergence** at rate $\mathcal{O}(1/t)$ under **strong convexity** with Lipschitz continuous $f, \nabla f, \nabla^2 f$ [Gürbüzbalaban et al., 2021].
 - SGD has min-max rate $\mathcal{O}(1/\sqrt{t})$ [Nemirovskij and Yudin, 1983].
- ▶ More results about **iteration complexity** are shown in [HaoChen and Sra, 2019; Nagaraj et al., 2019; Mishchenko et al., 2020, etc.]
- ▶ In nonconvex setting, $\liminf_{t \rightarrow \infty} \|\nabla f(x^t)\| = 0$ and **iteration complexity** results [Nguyen et al., 2020].



Overview

Random Reshuffling (RR)

Standard Kurdyka-Łojasiewicz (KL) Analysis Framework

A New Analysis Framework and Main Results

Conclusion



KL Inequality

KL Inequality

$f : \mathbb{R}^n \rightarrow \mathbb{R}$ is said to satisfy the **KL property** at $\bar{x} \in \mathbb{R}^n$ if

- ▶ \exists a desingularization function ϱ for all $x \in U := \text{neighborhood of } \bar{x}$

$$\varrho'(|f(x) - f(\bar{x})|) \cdot \|\nabla f(x)\| \geq 1. \quad (\text{KL inequality})$$

Remark:

- ▶ Popular choice: $\varrho(x) = cx^{1-\theta}$, i.e., the **Łojasiewicz inequality**



$$|f(x) - f(\bar{x})|^\theta \leq c \|\nabla f(x)\|, \quad c > 0.$$

Here, $\theta \in [0, 1)$ is called the **KL exponent**.

- ▶ Very mild and general [Attouch et al. (2013)].



Geometric Interpretation




$$\varrho'(f(x) - f(\bar{x})) \cdot \|\nabla f(x)\| \geq 1 \iff \psi'(x) \geq 1.$$

- ▶ ψ is **sharp** (absolute value like).
- ▶ KL inequality characterizes 'curvature' of f around \bar{x} .



Geometric Interpretation: KL exponent θ

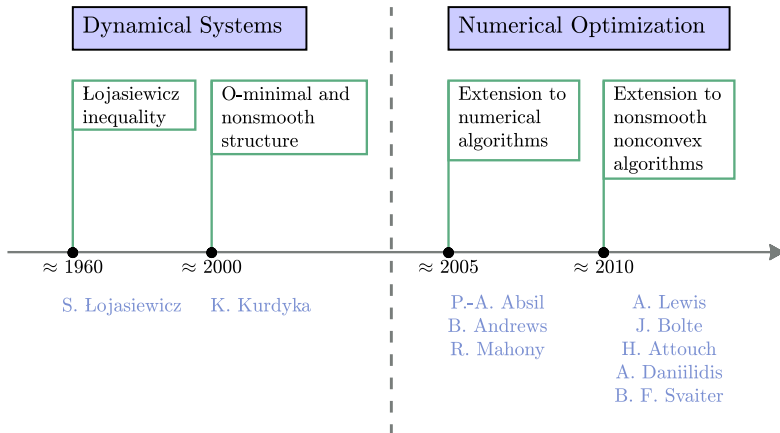


Functions with different KL exponent θ .

- ▶ Larger θ implies further from 'sharp' curvature.
- ▶ $\theta \in [0, \frac{1}{2}]$: good curvature.



KL History



Conditions


- **Sufficient decrease property:** (algorithmic)

$$f(x^{t+1}) \leq f(x^t) - c\|\nabla f(x^t)\|^2.$$

- This leads to weak convergence.

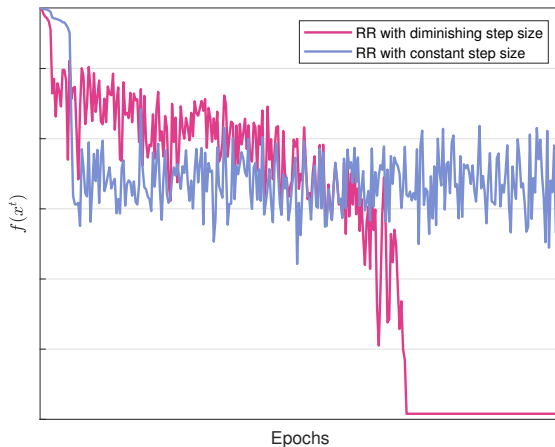
- f satisfies the **KL inequality** (problem intrinsic).

Proof Flow



The Fundamental Difference

- **Non-descent nature** and needs **diminishing step size**.



Overview

Random Reshuffling (RR)

Standard Kurdyka-Łojasiewicz (KL) Analysis Framework

A New Analysis Framework and Main Results

Conclusion



An Analysis Framework I

New Conditions (Algorithmic)

- **Approximate descent property:** $\exists b_1 > 0, b_2 > 0, \mu \geq 2$ s.t.

$$f(x^{t+1}) \leq f(x^t) - b_1 \alpha_t \|\nabla f(x^t)\|^2 + b_2 \alpha_t^\mu.$$

- **Relative error:** $\exists c_1 > 0, c_2 > 0, \nu \geq 2$ s.t.

$$\|x^{t+1} - x^t\| \leq c_1 \alpha_t \|\nabla f(x^t)\| + c_2 \alpha_t^\nu.$$

Observations:

- ' $b_2 \alpha_t^\mu$ ' is the **non-descent term**.
- RR only converges to a **neighborhood** if α_t is **constant step size**.



An Analysis Framework II

Diminishing step sizes (Algorithmic)

$$\alpha_t > 0, \quad \sum \alpha_t = \infty, \quad \text{and} \quad \sum \alpha_t^{\min\{\mu, \nu\}} < \infty.$$

- ▶ $\alpha_t = \mathcal{O}(\frac{1}{t^\gamma})$ with $\gamma = 1$ and $\frac{1}{2}$ are popular choices.

Consequence:

$$\text{Weak Convergence: } \lim_{t \rightarrow \infty} \|\nabla f(x^t)\| = 0.$$

Highlights:

- ▶ Approximate descent property + Diminishing step sizes \implies special ‘descent property’.
- ▶ Needs an elementary analysis (ε - δ arguments).




New Analysis Framework III

Additional Condition (Problem Intrinsic)

- **KL inequality** of f .

New Proof Flow



Highlights:

- A novel **descent-type condition** for the iterations.
- A variant of KL inequality (taking absolute value of " $f(x) - f(\bar{x})$ ").
- Combining **KL framework** with the dynamics of **diminishing step sizes**, etc.



Main Results

Assumptions

A.1 Each component function $f(\cdot, i)$ has **Lipschitz continuous gradient**.

A.2 Objective function f satisfies the **KL inequality**.

Thm: Convergence Results

Assume **A.1-A.2**. Using **diminishing step sizes** $\alpha_t = \mathcal{O}(1/t^\gamma)$, we obtain

► **strong (limit-point) convergence** $\lim_{t \rightarrow \infty} x^t = x^* \in \text{crit}(f)$.


► We have

$$\|x^t - x^*\| = \begin{cases} \mathcal{O}(1/t), & \text{if } \theta \in [0, 1/2], \\ \mathcal{O}(1/t^p), p \in (0, 1), & \text{if } \theta \in (1/2, 1). \end{cases}$$

Here, $\theta \in [0, 1)$ is the **KL exponent**.



Convergence Rate



Surface and Contour plot of the rates $\mathcal{O}(t^{-q})$ as a multifunction of the step size parameter $\alpha_t = \mathcal{O}(1/t^\gamma)$ with $\gamma \in (\frac{1}{2}, 1]$ and the KL exponent $\theta \in [0, 1)$.

- $\mathcal{O}(t^{-1})$ rate matches that of the the strongly convex setting.



Standing Out

	Nonconvex	Convex	Strongly Convex
Strong Convergence	<p>This paper</p>		<p>IG: Gürbüzbalaban et al. 19'</p> <p>RR: Gürbüzbalaban et al. 21'</p>
Weak Convergence	<p>IG:[bounded gradient] Luo & Tseng 94' Solodov & Zavriev 98' Tseng 98'</p> <p>RR: This paper</p>		
Iteration Complexity	<p>Nagaraj et al. 19' Mishchenko et al. 20' Nguyen et al. 21'</p>	<p>Shamir 16' Mishchenko et al. 20'</p>	<p>Nagaraj et al. 19' HaoChen & Sra 19' Mishchenko et al. 20' etc..</p>



Overview

Random Reshuffling (RR)

Standard Kurdyka-Łojasiewicz (KL) Analysis Framework

A New Analysis Framework and Main Results

Conclusion



Conclusion

- ▶ Random Reshuffling has been shown to have **strong limit-point convergence** in nonconvex setting under KL inequality.
- ▶ When KL exponent $\theta = 1/2$, we obtain rate $\|x^t - x^*\| = \mathcal{O}(1/t)$, which coincides with the rate under **strong convexity**.
- ▶ Significantly, our techniques/framework can potentially be utilized for a large class of **non-descent** algorithms with **diminishing step sizes**.



Reference I

- Absil, P.-A., R. Mahony, and B. Andrews
2005. Convergence of the iterates of descent methods for analytic cost functions. *SIAM Journal on Optimization*, 16(2):531–547.
- Attouch, H., J. Bolte, and B. F. Svaiter
2013. Convergence of descent methods for semi-algebraic and tame problems: proximal algorithms, forward-backward splitting, and regularized Gauss-Seidel methods. *Math. Program.*, 137(1-2, Ser. A):91–129.
- Gürbüzbalaban, M., A. Ozdaglar, and P. Parrilo
2021. Why random reshuffling beats stochastic gradient descent. *Mathematical Programming*, 186(1-2):49–84.
- Gürbüzbalaban, M., A. Ozdaglar, and P. A. Parrilo
2019. Convergence rate of incremental gradient and incremental Newton methods. *SIAM Journal on Optimization*, 29(4):2542–2565.
- HaoChen, J. Z. and S. Sra
2019. Random shuffling beats SGD after finite epochs. In *International Conference on Machine Learning*, Pp. 2624–2633.



Reference II

Luo, Z.-Q. and P. Tseng

1994. Analysis of an approximate gradient projection method with applications to the backpropagation algorithm. *Optimization Methods and Software*, 4(2):85–101.

Mishchenko, K., A. Khaled Ragab Bayoumi, and P. Richtárik

2020. Random reshuffling: Simple analysis with vast improvements. *Advances in Neural Information Processing Systems*, 33.

Nagaraj, D., P. Jain, and P. Netrapalli

2019. Sgd without replacement: Sharper rates for general smooth convex functions. In *International Conference on Machine Learning*, Pp. 4703–4711. PMLR.

Nemirovskij, A. S. and D. B. Yudin

1983. *Problem complexity and method efficiency in optimization*, Wiley-Interscience Series in Discrete Mathematics. Wiley, Chichester.

Nguyen, L. M., Q. Tran-Dinh, D. T. Phan, P. H. Nguyen, and M. van Dijk

2020. A unified convergence analysis for shuffling-type gradient methods. *arXiv preprint arXiv:2002.08246*.



Reference III

Solodov, M. V. and S. Zavriev

1998. Error stability properties of generalized gradient-type algorithms. *Journal of Optimization Theory and Applications*, 98(3):663–680.

Tseng, P.

1998. An incremental gradient(-projection) method with momentum term and adaptive stepsize rule. *SIAM Journal on Optimization*, 8(2):506–531.

